

Module 6 Assignment

kNN with bike frame geometries

General reminders about kNN:

- A type of pattern recognition algorithm
- A type of classification algorithm
- Non-parametric (that is, there's no output values like a coefficient in linear regression)
- Simple but has been successful in several areas, e.g. handwritten digits

In this homework, we're going to be using data on bicycles to try and predict which class a new bicycle fits into. We'll be using data on the geometry, or the shape, of bicycles. The below graphic labels the different parts of a bike's geometry. We don't need to know very much about bicycles to solve this problem, this is presented more for informational purposes.

All we need to know is that there are several different classes of bike, and each serves a different purpose. And each class of bike has a different geometry, or shape, and that this geometry can identify the class of a given bike. We'll use this information to build a classification model, specifically using kNN.



1. Imports

Run the imports cells to get started!

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

2. Read in the csv data

We have data on several measurements of bike geometries for 56cm sized bikes. We need to stick with the same size bike across all observations otherwise we'll conflate differences in scale.

Use `Pandas` and `read_csv` to read in the data. Save it as a DataFrame.

```
In [2]: # TODO read in the bike geo csv
# TODO read in the csv and recast the index to datetime format

bike_geometries = pd.read_csv('C:/Users/iamontes/Python Projects/Finance ML Certificat
print(bike_geometries.sample(10))
```

```
print(bike_geometries.shape)
```

	Brand	Year	Model	Size	Steer	Cat	Head Angle	\
65	Bianchi	1983.0	Randonneur	58.0		NaN	72.0	
140	Trek	1988.0	400	56.5		race	73.0	
305	Fuji	2006.0	Team Pro	56.0		race	73.5	
253	Miyata	1980.0	Gran Touring	58.0		NaN	73.0	
182	Trek	1977.0	TX900	56.0		race	73.0	
131	Mercian	1980.0	Unknown	61.0		race	73.0	
257	Gitane	1976.0	Tour de France	60.0		NaN	73.0	
50	Univega	1986.0	Gran Turismo	54.0		tour	72.0	
314	Cannondale	1995.0	R800	58.0		NaN	73.5	
123	Trek	2006.0	Pilot	56.0		NaN	72.5	

	Fork Offset	Seat Angle	Chain Stay	Wheelbase	Top Tube	BB Drop	Trail	\
65	5.20	72.0	43.0	103.5	57	6.6	55.8	
140	4.50	73.5	43.0	102.0	56.4	7.2	56.89	
305	4.50	73.0	41.0	98.6	56	6.9	53.78	
253	4.80	73.0	42.8	102.9	56.5	NaN	53.76	
182	4.50	73.0	42.0	100.3	56	6	56.89	
131	4.30	73.0	42.5	102.5	58	NaN	58.98	
257	6.35	73.0	41.9	101.6	57.8	NaN	37.55	
50	6.00	72.0	44.0	NaN	NaN	NaN	47.38	
314	4.12	74.0	40.6	101.0	57.5	NaN	57.74	
123	4.50	73.3	41.5	100.2	56.2	NaN	60.02	

	Flop
65	16.4
140	15.91
305	14.65
253	15.03
182	15.91
131	16.49
257	10.5
50	13.93
314	15.72
123	17.21

(356, 14)

3. Exploratory Data Analysis

We've been being very prescriptive in the assignments so far, but now that you have more experience, it's time to start exploring a little bit on your own! In the next couple of cells, explore the dataset we just imported.

Some things you'll want to consider:

- What kinds of columns are in the dataset
- What's the datatype for each column
- HINT: Be sure to print out the column names and watch out for extra whitespace! Look into the `strip` function to remove whitespace. Look into the `rename` function to rename a column
- How many of each class of bicycle are in our dataset? **HINT: the class is in the `Steer Cat` column**

```
In [3]: # TODO what columns are in the data
```

```
bike_geometries.columns
```

```
Out[3]: Index(['Brand', 'Year', 'Model', 'Size', 'Steer Cat', 'Head Angle',
          'Fork Offset', 'Seat Angle', 'Chain Stay', 'Wheelbase', 'Top Tube',
          'BB Drop', 'Trail ', 'Flop'],
          dtype='object')
```

```
In [4]: # TODO perhaps strip extra whitespace in the column names and rename the columns
bike_geometries.columns = bike_geometries.columns.str.strip()
```

```
print(bike_geometries.columns)
```

```
Index(['Brand', 'Year', 'Model', 'Size', 'Steer Cat', 'Head Angle',
          'Fork Offset', 'Seat Angle', 'Chain Stay', 'Wheelbase', 'Top Tube',
          'BB Drop', 'Trail', 'Flop'],
          dtype='object')
```

```
In [5]: # Replace 'old_name' and 'new_name' with actual column names
```

```
bike_geometries.rename(columns={
    'Trail ': 'Trail',
}, inplace=True)
```

```
In [6]: bike_geometries.columns
```

```
Out[6]: Index(['Brand', 'Year', 'Model', 'Size', 'Steer Cat', 'Head Angle',
          'Fork Offset', 'Seat Angle', 'Chain Stay', 'Wheelbase', 'Top Tube',
          'BB Drop', 'Trail', 'Flop'],
          dtype='object')
```

```
In [7]: # TODO use value_counts() to explore how many of each class of bike is in the dataset
bike_class_counts = bike_geometries['Steer Cat'].value_counts()
```

```
print(bike_class_counts)
```

```
Steer Cat
race      103
sport      49
cross      40
tour       18
crit        4
Name: count, dtype: int64
```

In [8]: *# TODO other EDA, such as .info() or .describe() or any other plot you think is necessary*

```
print(bike_geometries.info())

print(bike_geometries.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 356 entries, 0 to 355
Data columns (total 14 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Brand           356 non-null   object
 1   Year            345 non-null   float64
 2   Model           354 non-null   object
 3   Size            350 non-null   float64
 4   Steer Cat       214 non-null   object
 5   Head Angle      356 non-null   float64
 6   Fork Offset     356 non-null   float64
 7   Seat Angle      356 non-null   float64
 8   Chain Stay      354 non-null   float64
 9   Wheelbase       346 non-null   float64
10   Top Tube        338 non-null   object
11   BB Drop         186 non-null   object
12   Trail           349 non-null   object
13   Flop            349 non-null   object
dtypes: float64(7), object(7)
memory usage: 39.1+ KB
None
```

	Year	Size	Head Angle	Fork Offset	Seat Angle \
count	345.000000	350.000000	356.000000	356.000000	356.000000
mean	1987.866667	57.040571	72.835000	4.863287	73.021348
std	12.594511	1.576115	0.692122	0.668766	0.628564
min	1939.000000	52.000000	70.000000	2.700000	70.000000
25%	1978.000000	56.000000	72.500000	4.500000	73.000000
50%	1985.000000	57.000000	73.000000	4.500000	73.000000
75%	2004.000000	58.000000	73.000000	5.400000	73.300000
max	2007.000000	64.000000	74.500000	8.000000	75.300000

	Chain Stay	Wheelbase
count	354.000000	346.000000
mean	42.859887	104.309277
std	1.540584	48.134609
min	38.800000	96.800000
25%	41.500000	100.000000
50%	43.000000	101.300000
75%	44.000000	103.500000
max	47.000000	996.000000

4. Set up modeling dataset

Now it's time to get started with the modeling! The **first step is to create a modeling dataset**.

A modeling dataset includes only the columns that we'll want to use for modeling. You will want to **drop** columns that aren't useful. For example, you may want to (**hint**) drop columns that have a lot of **missing values**, such as NaN or nulls. Remember, our machine learning models don't work with null/NaN (Not a Number) inputs. So if a row has a NaN in even just one column, that whole row has to be thrown away by the model. But, instead of throwing out a whole row because of one NaN in a column, we can get rid of columns that have a lot of NaNs, and thereby save more rows for our model.

Your next steps:

- Find out which columns have a lot of missing data (NaN in the language of Pandas)
- Drop columns that have a lot of NaN s (but don't drop ALL columns with NaN s yet! Which ones might you want to keep?)
- You might want to keep some columns with NaN s, but then fill in the NaN with a value, e.g. 0
- Think about which columns you want to drop, and which ones with NaN s might benefit from replacing the NaN with a 0; you don't have to be a bike expert here, just take a guess and try several things
- Remember, to drop columns you can look into drop or you can subset your DataFrame by using a list of column names you want to keep, e.g.:

```
new_df = old_df[ ['column_a', 'column_b', 'column_f'] ]
```

- Drop rows with NaN in the Steer Cat column; remember, Steer Cat is our class, so we need to have a value there! Look into dropna()
- You may also want to drop any rows with a class type that's very uncommon in the data

In [9]: *# TODO Create a new dataframe that includes only our modeling columns*

```
nan_counts = bike_geometries.isna().sum()

print(nan_counts)

bike_geometries_new = bike_geometries.drop(columns=['BB Drop'])
```

```
Brand          0
Year           11
Model          2
Size           6
Steer Cat      142
Head Angle     0
Fork Offset    0
Seat Angle     0
Chain Stay     2
Wheelbase     10
Top Tube       18
BB Drop       170
Trail          7
Flop           7
dtype: int64
```

```
In [10]: #TODO what's the shape of your modeling dataset?
bike_geometries_new.shape
```

```
Out[10]: (356, 13)
```

```
In [11]: # TODO drop rows with NaN or NA for our output variable (Steer Cat); use dropna perhaps

bike_geometries_cleaned = bike_geometries_new.dropna(subset=['Steer Cat'])

# There are few observations on Year and Model, but I can't use median or frequency to
bike_geometries_cleaned = bike_geometries_cleaned.dropna(subset=['Year'])
bike_geometries_cleaned = bike_geometries_cleaned.dropna(subset=['Model'])

nan_counts = bike_geometries_cleaned.isna().sum()

print(nan_counts)
```

Brand	0
Year	0
Model	0
Size	2
Steer Cat	0
Head Angle	0
Fork Offset	0
Seat Angle	0
Chain Stay	1
Wheelbase	9
Top Tube	11
Trail	3
Flop	3

```
dtype: int64
```

```
In [12]: # TODO what's the shape of our modeling dataset now? Did we lose any observations?
print(bike_geometries.shape)

print(bike_geometries_cleaned.shape)
```

```
(356, 14)
(204, 13)
```

```
In [13]: 356-204
```

```
Out[13]: 152
```

- We lost 152 observations by dropping NaNs in Steer Cat, Year and Model. We interestingly didn't lose rows when dropping BB Drop which had 170 NaNs. Meaning that the NaNs for Steer Cat were also NaNs for BB Drop.
- I decided to drop Year and Model because there is not a normal distribution for those so I can input median. I can use frequency to input but I am pretty sure that it would be wrong.

```
In [14]: # TODO Look at the head() of your dataframe; what do you notice?
bike_geometries_cleaned.head()
```

Out[14]:

	Brand	Year	Model	Size	Steer Cat	Head Angle	Fork Offset	Seat Angle	Chain Stay	Wheelbase	Top Tube	Tra
12	Bruce Gordon	2002.0	Rock'nRoad	59.0	cross	72.0	5.0	73.0	45.0	107.5	58	57.
13	Specialized	1983.0	Expedition	58.0	cross	72.0	5.1	74.0	45.0	106.1	56.5	56.8
14	Bianchi	2003.0	San Remo	58.0	cross	72.0	5.0	73.0	44.0	104.7	57	Na
15	Miyata	1991.0	1000 LT	57.0	cross	72.0	5.0	72.0	45.0	104.7	56.5	57.
16	Nishiki	1974.0	International	58.0	cross	72.0	5.1	72.0	44.4	104.2	57.1	56.8

- There are still NAs in the dataset under Tail and Flop. We could replace with 0 but it's better to replace with the mean, median, or a regression that maps the missing value.
- I think median is best, not sensitive to outliers, and not complex like inputting a value via regression.

```
In [15]: # TODO Look into fillna to replace some missing values with, e.g., 0

# Specify columns to fill NaN values
columns_to_fill = ['Year', 'Wheelbase', 'Top Tube', 'Trail', 'Flop', 'Chain Stay', 'Si

# Fill NaN values with median for specified numeric columns
bike_geometries_cleaned[columns_to_fill] = bike_geometries_cleaned[columns_to_fill].fi

# Display cleaned DataFrame
print("\nDataFrame after filling NaN values with median:")
print(bike_geometries_cleaned)

nan_counts = bike_geometries_cleaned.isna().sum()

print(nan_counts)
```

DataFrame after filling NaN values with median:

	Brand	Year	Model	Size	Steer Cat	Head Angle \
12	Bruce Gordon	2002.0	Rock'nRoad	59.0	cross	72.0
13	Specialized	1983.0	Expedition	58.0	cross	72.0
14	Bianchi	2003.0	San Remo	58.0	cross	72.0
15	Miyata	1991.0	1000 LT	57.0	cross	72.0
16	Nishiki	1974.0	International	58.0	cross	72.0
..
310	Gaansari	2006.0	Van Cleve	56.0	race	73.5
311	Heron	2005.0	Rally/Road	56.0	race	73.5
323	Dawes	1978.0	Double Blue	58.0	crit	74.0
325	Holdsworth	1973.0	Racing Custom	61.0	crit	74.0
354	Waterford	2005.0	Track/Fixed-gear	56.0	crit	74.5

	Fork Offset	Seat Angle	Chain Stay	Wheelbase	Top Tube	Trail	Flop
12	5.00	73.0	45.0	107.5	58	57.9	17.02
13	5.10	74.0	45.0	106.1	56.5	56.85	16.71
14	5.00	73.0	44.0	104.7	57	55.87	15.71
15	5.00	72.0	45.0	104.7	56.5	57.9	17.02
16	5.10	72.0	44.4	104.2	57.1	56.85	16.71
..
310	4.25	72.5	42.0	100.3	57	56.39	15.36
311	4.25	72.5	42.5	100.3	57	56.39	15.36
323	3.80	73.0	43.2	101.3	59.7	57.96	15.36
325	3.18	75.0	41.3	99.1	56	64.41	17.07
354	3.00	75.0	39.5	97.0	56	63.16	16.26

[204 rows x 13 columns]

```
Brand      0
Year       0
Model      0
Size       0
Steer Cat  0
Head Angle 0
Fork Offset 0
Seat Angle 0
Chain Stay 0
Wheelbase  0
Top Tube   0
Trail      0
Flop       0
dtype: int64
```

In [16]: *# TODO Look at your dataset again with head() to see if the missing values were filled*
bike_geometries_cleaned.head()

Out[16]:

	Brand	Year	Model	Size	Steer Cat	Head Angle	Fork Offset	Seat Angle	Chain Stay	Wheelbase	Top Tube	Tra
12	Bruce Gordon	2002.0	Rock'nRoad	59.0	cross	72.0	5.0	73.0	45.0	107.5	58	57.
13	Specialized	1983.0	Expedition	58.0	cross	72.0	5.1	74.0	45.0	106.1	56.5	56.8
14	Bianchi	2003.0	San Remo	58.0	cross	72.0	5.0	73.0	44.0	104.7	57	55.8
15	Miyata	1991.0	1000 LT	57.0	cross	72.0	5.0	72.0	45.0	104.7	56.5	57.
16	Nishiki	1974.0	International	58.0	cross	72.0	5.1	72.0	44.4	104.2	57.1	56.8


```
In [17]: # TODO are there any classes (Steer Cat) that you want to drop from  
# your modeling dataset because there are very few of them? If so, do so here!  
bike_geometries_cleaned['Steer Cat'].value_counts().sum
```

```
Out[17]: <bound method NDFrame._add_numeric_operations.<locals>.sum of Steer Cat  
race      98  
sport     48  
cross     38  
tour      17  
crit       3  
Name: count, dtype: int64>
```

```
In [18]: # Drop rows where 'Steer Cat' is "crit"  
bike_geometries_cleaned = bike_geometries_cleaned[bike_geometries_cleaned['Steer Cat']
```

```
In [19]: # TODO what's the shape of your dataset now?  
bike_geometries_cleaned.shape
```

```
Out[19]: (201, 13)
```

```
In [20]: # TODO what's the value_counts() of your classes now?  
  
# Print Value counts after removing crit  
bike_geometries_cleaned['Steer Cat'].value_counts().sum
```

```
Out[20]: <bound method NDFrame._add_numeric_operations.<locals>.sum of Steer Cat  
race      98  
sport     48  
cross     38  
tour      17  
Name: count, dtype: int64>
```

```
In [21]: # Print Model  
print(bike_geometries_cleaned['Model'].value_counts().sum)  
# Print Brand  
print(bike_geometries_cleaned['Brand'].value_counts().sum)
```

```
<bound method NDFrame._add_numeric_operations.<locals>.sum of Model
930      9
510      7
730      6
Tour de France  6
Super Corsa    6
..
Accordo        1
Le Mans 12     1
Brava          1
Roubaix Pro    1
Rally/Road     1
Name: count, Length: 127, dtype: int64>
<bound method NDFrame._add_numeric_operations.<locals>.sum of Brand
Trek          70
Miyata        19
Gitane        13
Gaansari      10
Bianchi       8
Specialized   7
Bridgestone   7
Mercian       6
Ebisu         6
Jamis         5
Litespeed     4
Rivendell     4
Centurian     4
Fuji          4
LeMond        4
Schwinn       3
Nishiki       3
Lemond        2
Kogswell      2
Habanero      2
Velo Orange   1
Cinelli       1
Waterford     1
Holdsworth    1
Windsor       1
Merckx        1
Bruce Gordon  1
IF            1
Salsa         1
Soma          1
Terry         1
Univega       1
Victoria      1
Zabrakenko    1
Hetchins      1
Dawes         1
Ferrare       1
Heron         1
Name: count, dtype: int64>
```

```
In [22]: bike_geometries_cleaned = bike_geometries_cleaned.drop(columns=['Model'])
```

```
In [23]: # Count occurrences of each brand
brand_counts = bike_geometries_cleaned['Brand'].value_counts()
```

```
# Identify brands with fewer than 7 occurrences
brands_to_replace = brand_counts[brand_counts < 4].index

# Replace those brands with "Other"
bike_geometries_cleaned['Brand'] = bike_geometries_cleaned['Brand'].replace(brands_to_

# Display the modified DataFrame
print(bike_geometries_cleaned)
```

	Brand	Year	Size	Steer	Cat	Head	Angle	Fork	Offset	Seat	Angle	\
12	Other	2002.0	59.0		cross		72.0		5.00		73.0	
13	Specialized	1983.0	58.0		cross		72.0		5.10		74.0	
14	Bianchi	2003.0	58.0		cross		72.0		5.00		73.0	
15	Miyata	1991.0	57.0		cross		72.0		5.00		72.0	
16	Other	1974.0	58.0		cross		72.0		5.10		72.0	
..	
307	Other	2006.0	56.0		race		73.5		4.30		73.5	
308	Bianchi	2006.0	57.0		race		73.5		4.30		73.5	
309	Gaansari	2006.0	58.0		race		73.5		4.25		72.5	
310	Gaansari	2006.0	56.0		race		73.5		4.25		72.5	
311	Other	2005.0	56.0		race		73.5		4.25		72.5	

	Chain	Stay	Wheelbase	Top	Tube	Trail	Flop
12	45.0		107.5		58	57.9	17.02
13	45.0		106.1		56.5	56.85	16.71
14	44.0		104.7		57	55.87	15.71
15	45.0		104.7		56.5	57.9	17.02
16	44.4		104.2		57.1	56.85	16.71
..
307	40.4		98.1		56	55.87	15.21
308	40.6		100.3		56	55.87	15.21
309	42.0		100.3		58.5	56.39	15.36
310	42.0		100.3		57	56.39	15.36
311	42.5		100.3		57	56.39	15.36

[201 rows x 12 columns]

In [24]: `bike_geometries_cleaned['Brand'].value_counts().sum`

Out[24]: <bound method NDFrame._add_numeric_operations.<locals>.sum of Brand

Trek	70
Other	30
Miyata	19
Gitane	13
Gaansari	10
Bianchi	8
Specialized	7
Bridgestone	7
Ebisu	6
Mercian	6
Jamis	5
Rivendell	4
Litespeed	4
Centurian	4
Fuji	4
LeMond	4

Name: count, dtype: int64>

In [25]: `# Cast specified columns to float`
`bike_geometries_cleaned['Top Tube'] = bike_geometries_cleaned['Top Tube'].astype(float`

```
bike_geometries_cleaned['Trail'] = bike_geometries_cleaned['Trail'].astype(float)
bike_geometries_cleaned['Flop'] = bike_geometries_cleaned['Flop'].astype(float)
```

- Looks like Model has too much cardinality, Brand is the only one we can rescue if we transform to other.

```
In [26]: # Perform one-hot encoding on the 'Model' and 'Brand' columns
one_hot_encoded = pd.get_dummies(bike_geometries_cleaned[['Brand']], prefix=['Brand'])

# Convert True/False to 1/0
one_hot_encoded = one_hot_encoded.astype(int)

# Concatenate the one-hot encoded columns with the original DataFrame
bike_geometries_encoded = pd.concat([bike_geometries_cleaned, one_hot_encoded], axis=1)

# Optionally drop the original 'Model' and 'Brand' columns if no longer needed
bike_geometries_encoded.drop('Brand', axis=1, inplace=True)

# Display the DataFrame after one-hot encoding
print("\nDataFrame after one-hot encoding:")
print(bike_geometries_encoded)
```

DataFrame after one-hot encoding:

	Year	Size	Steer	Cat	Head	Angle	Fork	Offset	Seat	Angle	Chain	Stay	\
12	2002.0	59.0		cross		72.0		5.00		73.0		45.0	
13	1983.0	58.0		cross		72.0		5.10		74.0		45.0	
14	2003.0	58.0		cross		72.0		5.00		73.0		44.0	
15	1991.0	57.0		cross		72.0		5.00		72.0		45.0	
16	1974.0	58.0		cross		72.0		5.10		72.0		44.4	
..	
307	2006.0	56.0		race		73.5		4.30		73.5		40.4	
308	2006.0	57.0		race		73.5		4.30		73.5		40.6	
309	2006.0	58.0		race		73.5		4.25		72.5		42.0	
310	2006.0	56.0		race		73.5		4.25		72.5		42.0	
311	2005.0	56.0		race		73.5		4.25		72.5		42.5	

	Wheelbase	Top	Tube	Trail	...	Brand_Gitane	Brand_Jamis	Brand_LeMond	\
12	107.5		58.0	57.90	...	0	0	0	
13	106.1		56.5	56.85	...	0	0	0	
14	104.7		57.0	55.87	...	0	0	0	
15	104.7		56.5	57.90	...	0	0	0	
16	104.2		57.1	56.85	...	0	0	0	
..	
307	98.1		56.0	55.87	...	0	0	0	
308	100.3		56.0	55.87	...	0	0	0	
309	100.3		58.5	56.39	...	0	0	0	
310	100.3		57.0	56.39	...	0	0	0	
311	100.3		57.0	56.39	...	0	0	0	

	Brand_Litespeed	Brand_Mercian	Brand_Miyata	Brand_Other	\
12	0	0	0	1	
13	0	0	0	0	
14	0	0	0	0	
15	0	0	1	0	
16	0	0	0	1	
..	
307	0	0	0	1	
308	0	0	0	0	
309	0	0	0	0	
310	0	0	0	0	
311	0	0	0	1	

	Brand_Rivendell	Brand_Specialized	Brand_Trek
12	0	0	0
13	0	1	0
14	0	0	0
15	0	0	0
16	0	0	0
..
307	0	0	0
308	0	0	0
309	0	0	0
310	0	0	0
311	0	0	0

[201 rows x 27 columns]

5. kNN classification

Now that our modeling dataset is ready, we can start building our kNN model! Run the import cell below to get started. Notice we'll also be using a train/test split!

```
In [27]: # TODO run this cell
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
from sklearn.model_selection import train_test_split
```

Before we can run our model, we have to create our X matrix and y vector. To do this, you can take your modeling dataset from above and drop the dependent variable (`Steer Cat`) to create the X input matrix. Then take only `Steer Cat` to create your y vector.

```
In [28]: # TODO create X matrix and y vector from columns

X = bike_geometries_encoded.drop('Steer Cat', axis=1)
y = bike_geometries_encoded['Steer Cat']

print(X.shape , y.shape)
#''''''''
print(X.sample(5))
#''''''''
print(y.sample(5))
```

```
(201, 26) (201,)
      Year  Size  Head Angle  Fork Offset  Seat Angle  Chain Stay  Wheelbase  \
221  1980.0  56.0    73.0    5.5    73.0    44.5    103.5
190  2005.0  58.0    73.0    4.5    72.0    43.5    100.3
224  1980.0  56.0    73.0    5.5    73.0    44.5    103.5
217  1979.0  56.0    73.0    5.5    73.0    44.5    103.5
176  1988.0  56.0    73.0    4.5    74.0    40.8    99.2

      Top Tube  Trail  Flop  ...  Brand_Gitane  Brand_Jamis  Brand_LeMond  \
221      56.0  46.44  12.98  ...           0           0           0
190      57.0  56.89  15.91  ...           0           0           0
224      56.0  46.44  12.98  ...           0           0           0
217      56.0  46.44  12.98  ...           0           0           0
176      55.0  56.89  15.91  ...           0           0           0

      Brand_Litespeed  Brand_Mercian  Brand_Miyata  Brand_Other  \
221                  0              0            0            0
190                  0              0            0            1
224                  0              0            0            0
217                  0              0            0            0
176                  0              0            1            0

      Brand_Rivendell  Brand_Specialized  Brand_Trek
221                  0                  0            1
190                  0                  0            0
224                  0                  0            1
217                  0                  0            1
176                  0                  0            0
```

[5 rows x 26 columns]

173 race

138 race

311 race

42 tour

106 sport

Name: Steer Cat, dtype: object

```
In [29]: # TODO run info on X
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 201 entries, 12 to 311
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Year                  201 non-null   float64
1   Size                  201 non-null   float64
2   Head Angle            201 non-null   float64
3   Fork Offset           201 non-null   float64
4   Seat Angle            201 non-null   float64
5   Chain Stay            201 non-null   float64
6   Wheelbase             201 non-null   float64
7   Top Tube              201 non-null   float64
8   Trail                 201 non-null   float64
9   Flop                  201 non-null   float64
10  Brand_Bianchi          201 non-null   int32
11  Brand_Bridgestone      201 non-null   int32
12  Brand_Centurian        201 non-null   int32
13  Brand_Ebisu            201 non-null   int32
14  Brand_Fuji             201 non-null   int32
15  Brand_Gaansari         201 non-null   int32
16  Brand_Gitane           201 non-null   int32
17  Brand_Jamis            201 non-null   int32
18  Brand_LeMond           201 non-null   int32
19  Brand_Litespeed        201 non-null   int32
20  Brand_Mercian          201 non-null   int32
21  Brand_Miyata           201 non-null   int32
22  Brand_Other            201 non-null   int32
23  Brand_Rivendell        201 non-null   int32
24  Brand_Specialized      201 non-null   int32
25  Brand_Trek             201 non-null   int32
dtypes: float64(10), int32(16)
memory usage: 29.8 KB
```

```
In [30]: # TODO in info above, did you see any columns that should be floats but aren't?
# If so, cast them as floats now! e.g. X = X.astype(float)
# TODO run info on X
X.info()
```



```

<class 'pandas.core.frame.DataFrame'>
Index: 201 entries, 12 to 311
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Year                  201 non-null   float64
1   Size                  201 non-null   float64
2   Head Angle            201 non-null   float64
3   Fork Offset           201 non-null   float64
4   Seat Angle            201 non-null   float64
5   Chain Stay            201 non-null   float64
6   Wheelbase             201 non-null   float64
7   Top Tube              201 non-null   float64
8   Trail                 201 non-null   float64
9   Flop                  201 non-null   float64
10  Brand_Bianchi          201 non-null   int32
11  Brand_Bridgestone      201 non-null   int32
12  Brand_Centurian        201 non-null   int32
13  Brand_Ebisu            201 non-null   int32
14  Brand_Fuji             201 non-null   int32
15  Brand_Gaansari         201 non-null   int32
16  Brand_Gitane           201 non-null   int32
17  Brand_Jamis            201 non-null   int32
18  Brand_LeMond           201 non-null   int32
19  Brand_Litespeed        201 non-null   int32
20  Brand_Mercian          201 non-null   int32
21  Brand_Miyata           201 non-null   int32
22  Brand_Other            201 non-null   int32
23  Brand_Rivendell        201 non-null   int32
24  Brand_Specialized      201 non-null   int32
25  Brand_Trek             201 non-null   int32
dtypes: float64(10), int32(16)
memory usage: 29.8 KB

```

```

In [31]: # TODO check info again to make sure any casts worked
X.info()

```

```
<class 'pandas.core.frame.DataFrame'>
Index: 201 entries, 12 to 311
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Year                  201 non-null   float64
1   Size                  201 non-null   float64
2   Head Angle            201 non-null   float64
3   Fork Offset           201 non-null   float64
4   Seat Angle            201 non-null   float64
5   Chain Stay            201 non-null   float64
6   Wheelbase             201 non-null   float64
7   Top Tube              201 non-null   float64
8   Trail                 201 non-null   float64
9   Flop                  201 non-null   float64
10  Brand_Bianchi          201 non-null   int32
11  Brand_Bridgestone      201 non-null   int32
12  Brand_Centurian        201 non-null   int32
13  Brand_Ebisu            201 non-null   int32
14  Brand_Fuji             201 non-null   int32
15  Brand_Gaansari         201 non-null   int32
16  Brand_Gitane           201 non-null   int32
17  Brand_Jamis            201 non-null   int32
18  Brand_LeMond           201 non-null   int32
19  Brand_Litespeed        201 non-null   int32
20  Brand_Mercian          201 non-null   int32
21  Brand_Miyata           201 non-null   int32
22  Brand_Other            201 non-null   int32
23  Brand_Rivendell        201 non-null   int32
24  Brand_Specialized      201 non-null   int32
25  Brand_Trek             201 non-null   int32
dtypes: float64(10), int32(16)
memory usage: 29.8 KB
```

In [32]: `X.head()`

Out[32]:

	Year	Size	Head Angle	Fork Offset	Seat Angle	Chain Stay	Wheelbase	Top Tube	Trail	Flop	...	Brand_Gitane	Brar
12	2002.0	59.0	72.0	5.0	73.0	45.0	107.5	58.0	57.90	17.02	...	0	
13	1983.0	58.0	72.0	5.1	74.0	45.0	106.1	56.5	56.85	16.71	...	0	
14	2003.0	58.0	72.0	5.0	73.0	44.0	104.7	57.0	55.87	15.71	...	0	
15	1991.0	57.0	72.0	5.0	72.0	45.0	104.7	56.5	57.90	17.02	...	0	
16	1974.0	58.0	72.0	5.1	72.0	44.4	104.2	57.1	56.85	16.71	...	0	

5 rows × 26 columns

5a) kNN with train/test split

Before we build any model, ever, we have to do a train/test split. Go ahead and do a 70/30 train/test split now. Use the built-in sklearn method `train_test_split`.

Then fill in the blank (...) in the code below to get your first kNN model working!

```
In [33]: import math
print(X.shape)
result = math.sqrt(201)

print(result)
```

```
(201, 26)
14.177446878757825
```

```
In [34]: X
```

```
Out[34]:
```

	Year	Size	Head Angle	Fork Offset	Seat Angle	Chain Stay	Wheelbase	Top Tube	Trail	Flop	...	Brand_Gitane	Br
12	2002.0	59.0	72.0	5.00	73.0	45.0	107.5	58.0	57.90	17.02	...	0	
13	1983.0	58.0	72.0	5.10	74.0	45.0	106.1	56.5	56.85	16.71	...	0	
14	2003.0	58.0	72.0	5.00	73.0	44.0	104.7	57.0	55.87	15.71	...	0	
15	1991.0	57.0	72.0	5.00	72.0	45.0	104.7	56.5	57.90	17.02	...	0	
16	1974.0	58.0	72.0	5.10	72.0	44.4	104.2	57.1	56.85	16.71	...	0	
...	
307	2006.0	56.0	73.5	4.30	73.5	40.4	98.1	56.0	55.87	15.21	...	0	
308	2006.0	57.0	73.5	4.30	73.5	40.6	100.3	56.0	55.87	15.21	...	0	
309	2006.0	58.0	73.5	4.25	72.5	42.0	100.3	58.5	56.39	15.36	...	0	
310	2006.0	56.0	73.5	4.25	72.5	42.0	100.3	57.0	56.39	15.36	...	0	
311	2005.0	56.0	73.5	4.25	72.5	42.5	100.3	57.0	56.39	15.36	...	0	

201 rows × 26 columns

```
In [35]: y
```

```
Out[35]: 12    cross
13    cross
14    cross
15    cross
16    cross
...
307   race
308   race
309   race
310   race
311   race
Name: Steer Cat, Length: 201, dtype: object
```

```
In [36]: from sklearn.preprocessing import LabelEncoder

# Initialize the LabelEncoder
label_encoder = LabelEncoder()
```

```
# Fit and transform the labels
y_encoded = label_encoder.fit_transform(y)
```

In [37]: y

```
Out[37]: 12    cross
          13    cross
          14    cross
          15    cross
          16    cross
          ...
          307   race
          308   race
          309   race
          310   race
          311   race
          Name: Steer Cat, Length: 201, dtype: object
```

In [38]: y_encoded

```
Out[38]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 3, 3,
          3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
          2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
          2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
          2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 1, 1])
```

- 14 neighbors, seems like a lot, but lets test different amounts

```
In [39]: # TODO A 70/30 train/test split using sklearn's train_test_split

from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=
```

```
In [40]: # fit on train
knn = KNeighborsClassifier(n_neighbors=7)
knn.fit(X_train, y_train)
```

```
Out[40]: KNeighborsClassifier
KNeighborsClassifier(n_neighbors=7)
```

```
In [41]: print("X_train shape:", X_train.shape)
          print("y_train shape:", y_train.shape)
          print("\nData types:")
          print("X_train types:\n", X_train.dtypes)
          print("y_train type:", y_train.dtype)
```

```
X_train shape: (140, 26)
y_train shape: (140,)
```

Data types:

X_train types:

Year	float64
Size	float64
Head Angle	float64
Fork Offset	float64
Seat Angle	float64
Chain Stay	float64
Wheelbase	float64
Top Tube	float64
Trail	float64
Flop	float64
Brand_Bianchi	int32
Brand_Bridgestone	int32
Brand_Centurian	int32
Brand_Ebisu	int32
Brand_Fuji	int32
Brand_Gaansari	int32
Brand_Gitane	int32
Brand_Jamis	int32
Brand_LeMond	int32
Brand_Litespeed	int32
Brand_Mercian	int32
Brand_Miyata	int32
Brand_Other	int32
Brand_Rivendell	int32
Brand_Specialized	int32
Brand_Trek	int32
dtype:	object
y_train type:	object

```
In [42]: print("\nMissing values in X_train:\n", X_train.isnull().sum())
print("Missing values in y_train:\n", y_train.isnull().sum())
```

Missing values in X_train:

Year	0
Size	0
Head Angle	0
Fork Offset	0
Seat Angle	0
Chain Stay	0
Wheelbase	0
Top Tube	0
Trail	0
Flop	0
Brand_Bianchi	0
Brand_Bridgestone	0
Brand_Centurian	0
Brand_Ebisu	0
Brand_Fuji	0
Brand_Gaansari	0
Brand_Gitane	0
Brand_Jamis	0
Brand_LeMond	0
Brand_Litespeed	0
Brand_Mercian	0
Brand_Miyata	0
Brand_Other	0
Brand_Rivendell	0
Brand_Specialized	0
Brand_Trek	0

dtype: int64

Missing values in y_train:

0

```
In [43]: print("\nMissing values in X_train:\n", X_test.isnull().sum())  
         print("Missing values in y_train:\n", y_test.isnull().sum())
```

Missing values in X_train:

Year	0
Size	0
Head Angle	0
Fork Offset	0
Seat Angle	0
Chain Stay	0
Wheelbase	0
Top Tube	0
Trail	0
Flop	0
Brand_Bianchi	0
Brand_Bridgestone	0
Brand_Centurian	0
Brand_Ebisu	0
Brand_Fuji	0
Brand_Gaansari	0
Brand_Gitane	0
Brand_Jamis	0
Brand_LeMond	0
Brand_Litespeed	0
Brand_Mercian	0
Brand_Miyata	0
Brand_Other	0
Brand_Rivendell	0
Brand_Specialized	0
Brand_Trek	0

dtype: int64

Missing values in y_train:

0

```
In [44]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import LabelEncoder
from sklearn import metrics
import numpy as np
```

```
In [45]: print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("\nX_train data types:\n", X_train.dtypes)
print("\nX_test data types:\n", X_test.dtypes)
```

X_train shape: (140, 26)

X_test shape: (61, 26)

X_train data types:

Year	float64
Size	float64
Head Angle	float64
Fork Offset	float64
Seat Angle	float64
Chain Stay	float64
Wheelbase	float64
Top Tube	float64
Trail	float64
Flop	float64
Brand_Bianchi	int32
Brand_Bridgestone	int32
Brand_Centurian	int32
Brand_Ebisu	int32
Brand_Fuji	int32
Brand_Gaansari	int32
Brand_Gitane	int32
Brand_Jamis	int32
Brand_LeMond	int32
Brand_Litespeed	int32
Brand_Mercian	int32
Brand_Miyata	int32
Brand_Other	int32
Brand_Rivendell	int32
Brand_Specialized	int32
Brand_Trek	int32

dtype: object

X_test data types:

Year	float64
Size	float64
Head Angle	float64
Fork Offset	float64
Seat Angle	float64
Chain Stay	float64
Wheelbase	float64
Top Tube	float64
Trail	float64
Flop	float64
Brand_Bianchi	int32
Brand_Bridgestone	int32
Brand_Centurian	int32
Brand_Ebisu	int32
Brand_Fuji	int32
Brand_Gaansari	int32
Brand_Gitane	int32
Brand_Jamis	int32
Brand_LeMond	int32
Brand_Litespeed	int32
Brand_Mercian	int32
Brand_Miyata	int32
Brand_Other	int32
Brand_Rivendell	int32
Brand_Specialized	int32
Brand_Trek	int32

dtype: object


```
In [46]: # Convert your training and test sets to C-contiguous arrays
X_train = np.ascontiguousarray(X_train)
X_test = np.ascontiguousarray(X_test)

# Convert your training and test sets to C-contiguous arrays
y_train = np.ascontiguousarray(y_train)
y_test = np.ascontiguousarray(y_test)
```

```
In [47]: X_train
```

```
Out[47]: array([[1.978e+03, 5.600e+01, 7.300e+01, ..., 0.000e+00, 0.000e+00,
        1.000e+00],
        [1.984e+03, 5.800e+01, 7.200e+01, ..., 0.000e+00, 0.000e+00,
        0.000e+00],
        [1.984e+03, 5.700e+01, 7.200e+01, ..., 0.000e+00, 0.000e+00,
        0.000e+00],
        ...,
        [1.986e+03, 5.700e+01, 7.300e+01, ..., 0.000e+00, 0.000e+00,
        0.000e+00],
        [2.004e+03, 6.300e+01, 7.350e+01, ..., 0.000e+00, 0.000e+00,
        0.000e+00],
        [1.979e+03, 5.600e+01, 7.300e+01, ..., 0.000e+00, 0.000e+00,
        1.000e+00]])
```

```
In [48]: X_test
```

```
Out[48]: array([[1.980e+03, 5.600e+01, 7.300e+01, ..., 0.000e+00, 0.000e+00,
        1.000e+00],
        [1.988e+03, 5.600e+01, 7.200e+01, ..., 0.000e+00, 0.000e+00,
        0.000e+00],
        [1.984e+03, 5.800e+01, 7.200e+01, ..., 0.000e+00, 0.000e+00,
        0.000e+00],
        ...,
        [2.006e+03, 5.800e+01, 7.300e+01, ..., 0.000e+00, 0.000e+00,
        0.000e+00],
        [1.978e+03, 5.600e+01, 7.300e+01, ..., 0.000e+00, 0.000e+00,
        1.000e+00],
        [1.986e+03, 5.900e+01, 7.300e+01, ..., 0.000e+00, 0.000e+00,
        0.000e+00]])
```

```
In [49]: # Convert all columns to float if they are not already
#X_train = X_train.astype(float)
#X_test = X_test.astype(float)

# Convert all columns to float if they are not already
#y_train = y_train.astype(float)
#y_test = y_test.astype(float)
```

```
In [50]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=

# Fit the KNN classifier on the training data
knn = KNeighborsClassifier(n_neighbors=7)
knn.fit(X_train, y_train)

# Predict on test set
y_pred = knn.predict(X_test.values)

# Print accuracy metric
```

```
accuracy = metrics.accuracy_score(y_test, y_pred)
print("Accuracy of the KNN classifier:", accuracy)
```

C:\Users\iamontes\AppData\Local\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted with feature names

```
warnings.warn(
```

Accuracy of the KNN classifier: 0.8524590163934426

What's another metric we can use on a classification model, other than accuracy? A **confusion matrix**, perhaps? Use the built-in `sklearn` `confusion_matrix` from `sklearn.metrics`. Use the import below, then fill in the (...) to create a confusion matrix that compares the actual `y` against the predicted `y`.

```
In [51]: # TODO run the import
from sklearn.metrics import confusion_matrix
```

```
In [52]: # TODO fill in the blanks (...) and calculate the confusion matrix between the actual

# Calculate the confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Print the confusion matrix
print("Confusion Matrix:\n", cm)

# Normalize the confusion matrix to get percentages
cm_percentage = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis] * 100

# Print the confusion matrix as percentages with two decimal places
print("Confusion Matrix (Percentages):")
print(np.round(cm_percentage, 2)) # Round to two decimal places
```

Confusion Matrix:

```
[[ 4  4  0  0]
 [ 2 31  0  0]
 [ 0  1 15  1]
 [ 1  0  0  2]]
```

Confusion Matrix (Percentages):

```
[[50.  50.   0.   0. ]
 [ 6.06 93.94  0.   0. ]
 [ 0.   5.88 88.24  5.88]
 [33.33  0.   0.  66.67]]
```

5b) kNN with cross-validation: tuning for value of k

What value of `k` gives us the best kNN model? Let's use cross-validation to **tune the hyperparameter, k**. If you need a refresher on hyperparameter tuning with cross-validation, revisit the second instructor webinar. We'll walk you through it using kNN here.

In the webinar, we talked about how there are two uses for tuning a hyperparameter with cross-validation:

1. To choose the best value for a parameter (here, the `k` in kNN)

2. To choose the best model type between several algorithms (the next step below, choosing between kNN and logistic regression)

What are the inputs to `cross_val_score`? Use the [documentation](#) to fill in the (...) in the code skeleton below!

Remember, as we saw above, kNN is called `KNeighborsClassifier` in `sklearn`. If you need to, be sure to look at the [documentation](#).

```
In [53]: # TODO run the import for cross-val-score

from sklearn.model_selection import cross_val_score
```

```
In [54]: print("NaN values in X:", np.isnan(X).sum())
print("Infinite values in X:", np.isinf(X).sum())
```

```

NaN values in X: Year      0
Size                        0
Head Angle                 0
Fork Offset                0
Seat Angle                 0
Chain Stay                 0
Wheelbase                  0
Top Tube                   0
Trail                      0
Flop                       0
Brand_Bianchi              0
Brand_Bridgestone          0
Brand_Centurian            0
Brand_Ebisu                0
Brand_Fuji                 0
Brand_Gaansari             0
Brand_Gitane               0
Brand_Jamis                0
Brand_LeMond               0
Brand_Litespeed            0
Brand_Mercian              0
Brand_Miyata               0
Brand_Other                0
Brand_Rivendell            0
Brand_Specialized          0
Brand_Trek                 0
dtype: int64

```

```

Infinite values in X: Year      0
Size                        0
Head Angle                 0
Fork Offset                0
Seat Angle                 0
Chain Stay                 0
Wheelbase                  0
Top Tube                   0
Trail                      0
Flop                       0
Brand_Bianchi              0
Brand_Bridgestone          0
Brand_Centurian            0
Brand_Ebisu                0
Brand_Fuji                 0
Brand_Gaansari             0
Brand_Gitane               0
Brand_Jamis                0
Brand_LeMond               0
Brand_Litespeed            0
Brand_Mercian              0
Brand_Miyata               0
Brand_Other                0
Brand_Rivendell            0
Brand_Specialized          0
Brand_Trek                 0
dtype: int64

```

```
In [55]: X = np.ascontiguousarray(X.values.astype(float))
```

```
In [56]: # TODO (1) Run a 10-fold cross-validation with K=7 for kNN (the n_neighbors parameter)
#
```

```
# Create the KNN classifier with n_neighbors=7
knn = KNeighborsClassifier(n_neighbors=7)

# Run a 10-fold cross-validation
scores = cross_val_score(knn, X, y, cv=5, scoring='accuracy')

# Print the accuracy scores from cross-validation
print("Cross-validation accuracy scores:", scores)
```

Cross-validation accuracy scores: [0.70731707 0.775 0.9 0.775 0.8]

In [57]: *# TODO (2) Get average accuracy as an estimate of out-of-sample accuracy*

```
# Calculate average accuracy as an estimate of out-of-sample accuracy
average_accuracy = scores.mean()
print("Average accuracy (out-of-sample estimate):", average_accuracy)
```

Average accuracy (out-of-sample estimate): 0.7914634146341463

In [58]: *# TODO (3) Search for an optimal value of k from 1-30 (write a loop)*

```
# Initialize variables for k range and scores
k_range = list(range(1, 31))
k_scores = []

# Loop through k values from 1 to 30
for k in k_range:
    # Create a new KNeighborsClassifier with the current value of k
    knn = KNeighborsClassifier(n_neighbors=k)

    # Perform cross-validation and get accuracy scores
    scores = cross_val_score(knn, X, y, cv=10, scoring='accuracy')

    # Append the average score to k_scores
    k_scores.append(scores.mean())


# Print the average accuracy scores for each value of k with two decimal places
print("Average accuracy scores for each value of k:")
for k, score in zip(k_range, k_scores):
    print(f"k={k}: {score:.2f}")
```

Average accuracy scores for each value of k:

k=1: 0.84
k=2: 0.81
k=3: 0.84
k=4: 0.78
k=5: 0.79
k=6: 0.79
k=7: 0.81
k=8: 0.81
k=9: 0.80
k=10: 0.80
k=11: 0.79
k=12: 0.79
k=13: 0.78
k=14: 0.76
k=15: 0.77
k=16: 0.76
k=17: 0.74
k=18: 0.76
k=19: 0.76
k=20: 0.73
k=21: 0.74
k=22: 0.71
k=23: 0.74
k=24: 0.73
k=25: 0.74
k=26: 0.73
k=27: 0.73
k=28: 0.73
k=29: 0.72
k=30: 0.72

```
In [59]: # TODO plot the value of K for kNN (x-axis) versus the cross-validated accuracy (y-axis)
plt.figure(dpi=150)
plt.plot(k_range, k_scores, marker='o') # where are the scores for each value of k in
plt.xlabel('Value of K for KNN')
plt.ylabel('Cross-Validated Accuracy')
```

```
Out[59]: Text(0, 0.5, 'Cross-Validated Accuracy')
```



```

In [60]: # (4) TODO now try with k-fold = 3 (cv parameter)
# fill in the loop and replace the ...'s!

k_range = list(range(1, 100))
k_scores = []

for k in k_range:
    ## same code as above, just change cv=3
    # Create a new KNeighborsClassifier with the current value of k
    knn = KNeighborsClassifier(n_neighbors=k)

    # Perform cross-validation and get accuracy scores
    scores = cross_val_score(knn, X, y, cv=3, scoring='accuracy')
    k_scores.append(scores.mean())


plt.figure(dpi=150)
plt.plot(k_range, k_scores, marker='o')
plt.xlabel('Value of K for KNN')
plt.ylabel('Cross-Validated Accuracy')

```

```

Out[60]: Text(0, 0.5, 'Cross-Validated Accuracy')

```



5c) kNN with cross-validation: model selection of kNN vs. logistic

We know our optimal value of k , but let's use cross-validation to see how kNN compares against a logistic regression. This is called using cross-validation for **model selection**.

```
In [61]: # TODO Use 10-fold cross-validation with the best KNN model
# that is, set k equal to the k that gave you the best model above!

knn = KNeighborsClassifier(n_neighbors=3)
print(cross_val_score(knn, X, y, cv=10, scoring='accuracy').mean())

0.8416666666666666
```

We know the best value of k for a kNN, that's what we found in step 5b. We repeated that in the cell above so we can now compare it to a logistic regression. Use 10-fold cross-validation to see how well a logistic regression performs.

```
In [62]: # TODO run the import statement for a Logistic regression
from sklearn.linear_model import LogisticRegression
```

```
In [63]: # TODO run 10-fold cross-validation with Logistic regression

logreg = LogisticRegression(solver='liblinear', multi_class='auto') # use these values
print(cross_val_score(logreg, X, y, cv=10, scoring='accuracy').mean())
```


0.9452380952380952

6) Automating parameter tuning using GridSearchCV

We already used grid search above to find the best value of `k`. But there we used a manual loop that checked the values `k=1` to `k=30`. Is there an easier way? Of course there is! `sklearn` has a built-in `GridSearchCV` that combines a grid search and cross-validation. Import it and let's get to using it.

You may need to refer to the [documentation for GridSearchCV](#).

```
In [64]: # TODO run the import
from sklearn.model_selection import GridSearchCV
```

```
In [65]: # TODO define the parameter values that should be searched
# That is, what's the range of k's you want to try?

k_range = list(range(1, 31))
```

Now we just need to create the hyperparameter grid search input for `sklearn`. We need to pass it in as a dictionary, so let's create a dictionary now. We'll call our variable `param_grid` and it will hold the dictionary in this form:

```
{ name_of_hyperparameter : [ values, of, hyperparameter ] }
```

```
In [66]: # TODO create a parameter grid: map the parameter names to the values that should be s
# (just run this cell!)
param_grid = dict(n_neighbors=k_range)
print(param_grid)
```

```
{'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30]}
```

```
In [67]: # TODO instantiate the grid
# Use the param_grid from above and pass it into GridSearchCV
# Refer to the documentation if needed (Link above)

grid = GridSearchCV(KNeighborsClassifier(), param_grid, cv=10, scoring='accuracy')
```

```
In [68]: # TODO fit the grid with data
# Remember, we're fitting with our data (X, y)

grid.fit(X, y);
```

```
In [69]: # TODO examine the best model
# Just run this cell; you should have created grid in the cell above!

print(grid.best_score_)
print(grid.best_params_)
print(grid.best_estimator_)
```

```
0.8416666666666666
{'n_neighbors': 3}
KNeighborsClassifier(n_neighbors=3)
```

7) Reducing computational expense using RandomizedSearchCV

Instead of searching the entire search space (every possible combination of our hyperparameters), we can use a randomized search. A randomized grid search might not get us the **globally best model** but it will get us close enough! Refer to the second instructor webinar for more information on the difference between Grid Search and Randomized Grid Search.

```
In [70]: # TODO run this cell for importing randomized grid search
from sklearn.model_selection import RandomizedSearchCV
```

For the grid search, we used a parameter *grid* because we wanted to search every possible value of *k*. But here, instead of a *param_grid*, we'll be using a *param_dist* (parameter *distribution*) to sample from.

What's happening here is that we're not going to try every single value from our range of *k*. So let's say our range of *k* is 1-30. In grid search, we try 1, 2, 3, 4, ..., 30. In randomized grid search, we'll use a distribution to *sample* values of *k* from 1-30.

Here we'll use a **uniform distribution**. A uniform distribution will make every value of *k*, from 1 to 30, equally likely to be pulled. To tell this to *sklearn*, we'll be using the input *weights=* in *RandomizedSearchCV*. We'll first create a helper variable, though, called *param_dist* that will create a dictionary that we'll pass in to randomized search.

We don't *have to* create this helper variable, but it's just easier to create a dictionary with all of our hyperparameters and pass just one variable into *RandomizedSearchCV*. In this case, our two hyperparameters are *n_neighbors* and *weights*.

```
In [71]: # TODO specify "parameter distributions" rather than a "parameter grid"


k_range = list(range(1, 30))

param_dist = {
    'n_neighbors': k_range,
    'weights': ['uniform', 'distance']
}
```

```
In [72]: # n_iter controls the number of searches -- that is, how much of the potential grid do
# n_iter default = 10
# TODO try running with different values of n_iter!

rand = RandomizedSearchCV(knn, param_dist, cv=10, scoring='accuracy', \
                          n_iter=20, random_state=42)
rand.fit(X, y)
```

Out[72]:

In [73]: *# TODO examine the best model from our random grid search*

```
print(rand.best_score_)
print(rand.best_params_)
```

```
0.8766666666666666
{'weights': 'distance', 'n_neighbors': 3}
```

How well does a randomized search actually do? We can try running the randomized search multiple times and see how the output differs. If the output is similar in each run of the random search, we can be fairly confident that the random search will get us close enough to a full grid search.

In [75]: *# TODO run RandomizedSearchCV 20 times (with n_iter=10) and record the best score*

```
best_scores = []
knn = KNeighborsClassifier()

for _ in range(20):
    ##### your code here!
    # create a randomized grid search
    rand = RandomizedSearchCV(knn, param_distributions=param_dist, cv=10, scoring='acc
    # fit the rand search
    rand.fit(X, y)
    # append the best score from fit to best_scores
    rand = RandomizedSearchCV(knn, param_dist, cv=10, scoring='accuracy', n_iter=10)
    rand.fit(X, y)
    best_scores.append(round(rand.best_score_, 3))
print(best_scores)
```

```
[0.841, 0.847, 0.877, 0.846, 0.877, 0.856, 0.847, 0.856, 0.856, 0.872, 0.872, 0.877,
0.877, 0.877, 0.856, 0.856, 0.856, 0.856, 0.846, 0.877]
```

- I am confident that the random search will get me close enough to a full grid search because I am getting values between 85% and 88%.